

assignment_5

May 4, 2025

1 STA-6543 Assignment 5

Jason Gillette

1.1 Question 2.

For parts (a) through (c), indicate which of i. through iv. is correct. Justify your answer.

1.2 Question 2a.

The lasso, relative to least squares, is:

iii. Less flexible and hence will give improved prediction accuracy when its increase in bias is

Lasso works when the benefit of reduced variance outweighs the cost of increased bias. It

1.3 Question 2b.

Repeat (a) for ridge regression relative to least squares.

(iii) Less flexible and hence will give improved prediction accuracy when its increase in bias is

Least squares regression places no restrictions on a variable's coefficient to influence f

1.4 Question 2c.

Repeat (a) for non-linear methods relative to least squares.

(i) More flexible and hence will give improved prediction accuracy when its increase in bias is

Non-linear methods are not bound by a straight line and fit more complex patterns in the d

1.5 Question 9.

In this exercise, we will predict the number of applications received using the other variables in the College data set.

```
[27]: from ISLP import load_data
import pandas as pd

# Load College dataset
```

```
college = load_data('College')

# Look at initial structure
college.info()
college.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 777 entries, 0 to 776
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Private                777 non-null   category
1   Apps                   777 non-null   int64
2   Accept                 777 non-null   int64
3   Enroll                 777 non-null   int64
4   Top10perc              777 non-null   int64
5   Top25perc              777 non-null   int64
6   F.Undergrad            777 non-null   int64
7   P.Undergrad            777 non-null   int64
8   Outstate               777 non-null   int64
9   Room.Board             777 non-null   int64
10  Books                  777 non-null   int64
11  Personal               777 non-null   int64
12  PhD                   777 non-null   int64
13  Terminal               777 non-null   int64
14  S.F.Ratio              777 non-null   float64
15  perc.alumni            777 non-null   int64
16  Expend                 777 non-null   int64
17  Grad.Rate              777 non-null   int64
dtypes: category(1), float64(1), int64(16)
memory usage: 104.2 KB
```

```
[27]: Private Apps Accept Enroll Top10perc Top25perc F.Undergrad \
0      Yes 1660 1232 721 23 52 2885
1      Yes 2186 1924 512 16 29 2683
2      Yes 1428 1097 336 22 50 1036
3      Yes 417 349 137 60 89 510
4      Yes 193 146 55 16 44 249

      P.Undergrad Outstate Room.Board Books Personal PhD Terminal \
0           537 7440 3300 450 2200 70 78
1          1227 12280 6450 750 1500 29 30
2           99 11250 3750 400 1165 53 66
3           63 12960 5450 450 875 92 97
4          869 7560 4120 800 1500 76 72

      S.F.Ratio perc.alumni Expend Grad.Rate
```

0	18.1	12	7041	60
1	12.2	16	10527	56
2	12.9	30	8735	54
3	7.7	37	19016	59
4	11.9	2	10922	15

[28]: *### adding a cleaning step due to all the value errors downstream.*

```
# Rename columns to replace '.' with '_' for sml syntax
college.columns = [col.replace('.', '_') for col in college.columns]

# Convert 'Private' from Yes/No to 1/0 int
college['Private'] = college['Private'].cat.codes

# Drop any rows with NaNs
college = college.dropna()

# confirm casting and rename
college.info()
college.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 777 entries, 0 to 776
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Private                777 non-null   int8
1   Apps                   777 non-null   int64
2   Accept                 777 non-null   int64
3   Enroll                 777 non-null   int64
4   Top10perc              777 non-null   int64
5   Top25perc              777 non-null   int64
6   F_Undergrad            777 non-null   int64
7   P_Undergrad            777 non-null   int64
8   Outstate               777 non-null   int64
9   Room_Board            777 non-null   int64
10  Books                  777 non-null   int64
11  Personal               777 non-null   int64
12  PhD                    777 non-null   int64
13  Terminal               777 non-null   int64
14  S_F_Ratio              777 non-null   float64
15  perc_alumni            777 non-null   int64
16  Expend                 777 non-null   int64
17  Grad_Rate              777 non-null   int64
dtypes: float64(1), int64(16), int8(1)
memory usage: 104.1 KB
```

```
[28]:
```

	Private	Apps	Accept	Enroll	Top10perc	Top25perc	F_Undergrad	\
0	1	1660	1232	721	23	52	2885	
1	1	2186	1924	512	16	29	2683	
2	1	1428	1097	336	22	50	1036	
3	1	417	349	137	60	89	510	
4	1	193	146	55	16	44	249	

	P_Undergrad	Outstate	Room_Board	Books	Personal	PhD	Terminal	\
0	537	7440	3300	450	2200	70	78	
1	1227	12280	6450	750	1500	29	30	
2	99	11250	3750	400	1165	53	66	
3	63	12960	5450	450	875	92	97	
4	869	7560	4120	800	1500	76	72	

	S_F_Ratio	perc_alumni	Expend	Grad_Rate
0	18.1	12	7041	60
1	12.2	16	10527	56
2	12.9	30	8735	54
3	7.7	37	19016	59
4	11.9	2	10922	15

1.6 Question 9a.

Split the data set into a training set and a test set.

```
[29]: from sklearn.model_selection import train_test_split

# set random seed for reproducibility
RANDOM_STATE = 42

# Split the College data into features (X) and target (y)
X = college.drop(columns='Apps') # predictors
y = college['Apps']              # target: number of applications

# Perform train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=RANDOM_STATE
)

# sizes
print(f"Training set: {X_train.shape}, Test set: {X_test.shape}")
```

Training set: (543, 17), Test set: (234, 17)

1.7 Question 9b.

Fit a linear model using least squares on the training set, and report the test error obtained.

```
[30]: import pandas as pd
import statsmodels.formula.api as smf
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

# join X and y back into single training/test DataFrame
train = X_train.copy()
train['Apps'] = y_train

test = X_test.copy()
test['Apps'] = y_test

# Fit model
# Cite OpenAi o4 here, examples in book caused syntax error.
model = smf.ols(
    'Apps ~ Private + Accept + Enroll + Top10perc + Top25perc + F_Undergrad +
    P_Undergrad + Outstate + Room_Board + Books + Personal + PhD + Terminal +
    S_F_Ratio + perc_alumni + Expend + Grad_Rate',
    data=train
).fit()

# Predict
y_pred = model.predict(test)

# Test MSE
mse_test = mean_squared_error(test['Apps'], y_pred)
print(f"Test MSE: {mse_test:.4f}")
```

Test MSE: 1931803.1942

1.8 Question 9c.

Fit a ridge regression model on the training set, with α chosen by cross-validation. Report the test error obtained.

```
[31]: from sklearn.linear_model import RidgeCV
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
import numpy as np

# Ridge requires scaled predictors
ridge_model = make_pipeline(
    StandardScaler(),
    RidgeCV(alphas=10**np.linspace(10, -2, 100), store_cv_results=True)
)

# Fit ridge on training data
ridge_model.fit(X_train, y_train)
```

```

# Predict on test data
y_pred_ridge = ridge_model.predict(X_test)

# Test MSE
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
print(f"Test MSE (Ridge): {mse_ridge:.4f}")

# Best lambda chosen
best_lambda = ridge_model.named_steps['ridgecv'].alpha_
print(f"Optimal lambda (alpha) via CV: {best_lambda:.4f}")

```

Test MSE (Ridge): 1931467.5647
Optimal lambda (alpha) via CV: 0.0100

Ridge regression slightly improved test error by ~335 units, a modest gain, suggesting the least squares model may not be severely overfitting and Ridge Regression helped, but only marginally.

1.9 Question 9d.

Fit a lasso model on the training set, with λ chosen by cross-validation. Report the test error obtained, along with the number of non-zero coefficient estimates.

```

[32]: from sklearn.linear_model import LassoCV
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
import numpy as np
from sklearn.metrics import mean_squared_error

# Lasso model pipeline: standardize then cross-validated lasso
lasso_model = make_pipeline(
    StandardScaler(),
    LassoCV(alphas=10**np.linspace(10, -2, 100), cv=10, max_iter=10000,
    random_state=42)
)

# Fit model on training data
lasso_model.fit(X_train, y_train)

# Predict on test set
y_pred_lasso = lasso_model.predict(X_test)

# Compute test MSE
mse_lasso = mean_squared_error(y_test, y_pred_lasso)

# Extract best lambda and coefficients
lasso_cv = lasso_model.named_steps['lassocv']
best_lambda = lasso_cv.alpha_

```

```

nonzero_coefs = np.sum(lasso_cv.coef_ != 0)

print(f"Test MSE (Lasso): {mse_lasso:.4f}")
print(f"Optimal lambda (alpha): {best_lambda:.4f}")
print(f"Number of non-zero coefficients: {nonzero_coefs}")

```

```

Test MSE (Lasso): 1931787.4564
Optimal lambda (alpha): 0.0100
Number of non-zero coefficients: 17

```

Lasso chose the same number of predictors (17), meaning none were shrunk to zero; suggests all predictors are at least marginally useful. The MSE is slightly better than least squares, but slightly worse than ridge regression.

1.10 Question 9e.

Fit a PCR model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```

[35]: from sklearn.decomposition import PCA
      from sklearn.linear_model import LinearRegression
      from sklearn.preprocessing import StandardScaler
      from sklearn.pipeline import Pipeline
      from sklearn.model_selection import cross_val_score
      import numpy as np

      # total number of features (17)
      m_values = list(range(1, X_train.shape[1] + 1))
      cv_errors = []

      for m in m_values:
          # Create PCR pipeline
          pcr_pipeline = Pipeline([
              ('scale', StandardScaler()),
              ('pca', PCA(n_components=m)),
              ('linreg', LinearRegression())
          ])

          # Negative MSE (so higher is better) across 10 folds
          scores = cross_val_score(pcr_pipeline, X_train, y_train,
                                   scoring='neg_mean_squared_error', cv=10)

          # Average MSE (flip sign)
          cv_errors.append(-scores.mean())

      # Get best M
      best_m = m_values[np.argmin(cv_errors)]
      print(f"Best number of components (M): {best_m}")

```

```

# Final PCR model with best M
final_pcr = Pipeline([
    ('scale', StandardScaler()),
    ('pca', PCA(n_components=best_m)),
    ('linreg', LinearRegression())
])

final_pcr.fit(X_train, y_train)

# Predict on test set
y_pred_pcr = final_pcr.predict(X_test)

# Test MSE
from sklearn.metrics import mean_squared_error
mse_pcr = mean_squared_error(y_test, y_pred_pcr)
print(f"Test MSE (PCR): {mse_pcr:.4f}")

```

Best number of components (M): 17
Test MSE (PCR): 1931803.1942

The exact same test MSE as standard least squares... suggesting all variables are uniquely contributing and consistent with other model results?

1.11 Question 9f.

Fit a PLS model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation

```

[36]: from sklearn.cross_decomposition import PLSRegression
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error
import numpy as np

m_values = list(range(1, X_train.shape[1] + 1)) # Try 1 to 17 components
cv_errors = []

for m in m_values:
    # Build PLS pipeline
    pls_pipeline = Pipeline([
        ('scale', StandardScaler()),
        ('pls', PLSRegression(n_components=m))
    ])

    # Use negative MSE scoring, average over 10-fold CV
    scores = cross_val_score(pls_pipeline, X_train, y_train,

```



```

        scoring='neg_mean_squared_error', cv=10)

    cv_errors.append(-scores.mean())  # Flip sign for MSE

# Best number of components
best_m = m_values[np.argmin(cv_errors)]
print(f"Best number of components (M): {best_m}")

# Final PLS model with best M
final_pls = Pipeline([
    ('scale', StandardScaler()),
    ('pls', PLSRegression(n_components=best_m))
])

final_pls.fit(X_train, y_train)

# Predict on test set
y_pred_pls = final_pls.predict(X_test)

# Evaluate
mse_pls = mean_squared_error(y_test, y_pred_pls)
print(f"Test MSE (PLS): {mse_pls:.4f}")

```

Best number of components (M): 14

Test MSE (PLS): 1930317.7199

PLS performed best among all models so far, reducing test MSE by about 1,485. It only 14 used components meaning it eliminated 3 components that added noise rather than predictive signal. Since it uses y during component extraction, it focused only on signal relevant to prediction.

1.12 Question 9g.

Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these five approaches?

Model	Test MSE	Comments
Linear (OLS)	1,931,803.19	Baseline
Ridge (CV)	1,931,467.56	Slight improvement
Lasso (CV)	1,931,787.46	No real feature selection
PCR (CV)	1,931,803.19	Identical to OLS (M=17)
PLS (CV)	1,930,317.72	Best, using 14 components

The differences in test error are minimal, less than 0.1% between worst and best models. This suggests all 17 predictors are contributing at least modestly and regularization and dimensionality reduction helped only slightly

1.13 Question 11.

We will now try to predict per capita crime rate in the Boston dataset.

```
[37]: from ISLP import load_data
import pandas as pd

# Load the Boston dataset
boston = load_data('Boston')

# Preview the data
boston.info()
boston.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 506 entries, 0 to 505
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	crim	506 non-null	float64
1	zn	506 non-null	float64
2	indus	506 non-null	float64
3	chas	506 non-null	int64
4	nox	506 non-null	float64
5	rm	506 non-null	float64
6	age	506 non-null	float64
7	dis	506 non-null	float64
8	rad	506 non-null	int64
9	tax	506 non-null	int64
10	ptratio	506 non-null	float64
11	lstat	506 non-null	float64
12	medv	506 non-null	float64

```
dtypes: float64(10), int64(3)
```

```
memory usage: 51.5 KB
```

```
[37]:      crim    zn  indus  chas    nox    rm    age    dis  rad  tax  ptratio  \
0  0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900    1  296    15.3
1  0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671    2  242    17.8
2  0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671    2  242    17.8
3  0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622    3  222    18.7
4  0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622    3  222    18.7

      lstat  medv
0    4.98  24.0
1    9.14  21.6
2    4.03  34.7
3    2.94  33.4
4    5.33  36.2
```

1.14 Question 11a.

Try out some of the regression methods explored in this chapter, such as best subset selection, the lasso, ridge regression, and PCR. Present and discuss results for the approaches that you consider.

```
[44]: from sklearn.model_selection import train_test_split

from sklearn.linear_model import RidgeCV
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
import numpy as np

from sklearn.cross_decomposition import PLSRegression
from sklearn.model_selection import cross_val_score
from sklearn.pipeline import Pipeline

# Define predictors and response
X = boston.drop(columns='crim')
y = boston['crim']

# Train-test split (70/30)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=1
)

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# Fit and predict least squares
lm = LinearRegression()
lm.fit(X_train, y_train)
y_pred_lm = lm.predict(X_test)

# Evaluate least squares
mse_lm = mean_squared_error(y_test, y_pred_lm)
print(f"Test MSE (Least Squares): {mse_lm:.4f}")

# try ridge regression
ridge_model = make_pipeline(
    StandardScaler(),
    RidgeCV(alphas=10**np.linspace(10, -2, 100), store_cv_results=True)
)

ridge_model.fit(X_train, y_train)
y_pred_ridge = ridge_model.predict(X_test)

mse_ridge = mean_squared_error(y_test, y_pred_ridge)
best_lambda_ridge = ridge_model.named_steps['ridgecv'].alpha_
```

```

print(f"\nTest MSE (Ridge): {mse_ridge:.4f}")
print(f"Optimal lambda (Ridge): {best_lambda_ridge:.4f}")

# Try PLS
m_values = list(range(1, X_train.shape[1] + 1))
cv_errors = []

for m in m_values:
    pls_pipeline = Pipeline([
        ('scale', StandardScaler()),
        ('pls', PLSRegression(n_components=m))
    ])
    scores = cross_val_score(pls_pipeline, X_train, y_train,
                             scoring='neg_mean_squared_error', cv=10)
    cv_errors.append(-scores.mean())

# Choose best M
best_m = m_values[np.argmin(cv_errors)]
print(f"\nBest number of components (PLS): {best_m}")

# Fit final model
final_pls = Pipeline([
    ('scale', StandardScaler()),
    ('pls', PLSRegression(n_components=best_m))
])
final_pls.fit(X_train, y_train)
y_pred_pls = final_pls.predict(X_test)
mse_pls = mean_squared_error(y_test, y_pred_pls)

print(f"Test MSE (PLS): {mse_pls:.4f}")

```

Test MSE (Least Squares): 50.5105

Test MSE (Ridge): 50.3006

Optimal lambda (Ridge): 4.6416

Best number of components (PLS): 9

Test MSE (PLS): 50.5177

All three models predict per capita crime rate with similar accuracy. Ridge performed marginally better, suggesting that slight coefficient limiting helps stabilize the model. PLS didn't outperform despite reducing dimensionality, likely because most predictors carry some useful signal.

1.15 Question 11b.

Propose a model (or set of models) that seem to perform well on this data set, and justify your answer. Make sure that you are evaluating model performance using validation set error, cross-validation, or some other reasonable alternative, as opposed to using training error.

```
[45]: # Let's propose PCR and if poor performance, we'll go with ridge regression

from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error
import numpy as np

m_values = list(range(1, X_train.shape[1] + 1)) # Try 1 to 13 components
cv_errors = []

for m in m_values:
    pcr_pipeline = Pipeline([
        ('scale', StandardScaler()),
        ('pca', PCA(n_components=m)),
        ('linreg', LinearRegression())
    ])

    scores = cross_val_score(pcr_pipeline, X_train, y_train,
                             scoring='neg_mean_squared_error', cv=10)
    cv_errors.append(-scores.mean())

best_m = m_values[np.argmin(cv_errors)]
print(f"Best number of components (PCR): {best_m}")

# Final model with best M
final_pcr = Pipeline([
    ('scale', StandardScaler()),
    ('pca', PCA(n_components=best_m)),
    ('linreg', LinearRegression())
])

final_pcr.fit(X_train, y_train)
y_pred_pcr = final_pcr.predict(X_test)
mse_pcr = mean_squared_error(y_test, y_pred_pcr)

print(f"Test MSE (PCR): {mse_pcr:.4f}")
```

Best number of components (PCR): 12

Test MSE (PCR): 50.5105

Even though PCR used one fewer component than the full 13, the last principal component (PC13)

must have contributed very little predictive value, so the regression performance was effectively the same as OLS with all 13 original predictors.

Proposed Model is Ridge Regression Among the regression approaches tested, ridge regression produced the lowest test mean squared error (MSE), making it the strongest candidate for modeling per capita crime rate in the Boston dataset.

1.16 Question 11c.

Does your chosen model involve all of the features in the dataset? Why or why not?

Ridge regression includes all features in the model. While it penalizes large coefficients via regularization, it does not eliminate predictors altogether like other methods such as lasso. This is useful when all features are believed to carry some predictive information, which seems to be the case here.