

# assignment\_6

May 14, 2025

## 1 STA-6543 Assignment 6

Jason Gillette

### 1.1 Question 6

In this exercise, you will further analyze the Wage data set considered throughout this chapter.

```
[3]: import pandas as pd
      from ISLP import load_data

      # Load the Wage dataset
      Wage = load_data('Wage')

      # Show the first few rows
      Wage.head()
```

```
[3]:   year  age  maritl  race  education  region \
0  2006   18  1. Never Married  1. White  1. < HS Grad  2. Middle Atlantic
1  2004   24  1. Never Married  1. White  4. College Grad  2. Middle Atlantic
2  2003   45    2. Married  1. White  3. Some College  2. Middle Atlantic
3  2003   43    2. Married  3. Asian  4. College Grad  2. Middle Atlantic
4  2005   50    4. Divorced  1. White    2. HS Grad  2. Middle Atlantic

      jobclass  health health_ins  logwage  wage
0  1. Industrial    1. <=Good    2. No  4.318063  75.043154
1  2. Information  2. >=Very Good    2. No  4.255273  70.476020
2  1. Industrial    1. <=Good    1. Yes  4.875061  130.982177
3  2. Information  2. >=Very Good    1. Yes  5.041393  154.685293
4  2. Information    1. <=Good    1. Yes  4.318063  75.043154
```

### 1.2 Question 6a.

Perform polynomial regression to predict wage using age. Use cross-validation to select the optimal degree  $d$  for the polynomial. What degree was chosen, and how does this compare to the results of hypothesis testing using ANOVA? Make a plot of the resulting polynomial fit to the data.

```
[4]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score

# Define variables
X = Wage[['age']].values
y = Wage['wage'].values

# Try polynomial degrees from 1 to 10
degrees = range(1, 11)
cv_errors = []


for d in degrees:
    poly = PolynomialFeatures(degree=d, include_bias=False)
    X_poly = poly.fit_transform(X)
    model = LinearRegression()
    mse = -cross_val_score(model, X_poly, y, scoring='neg_mean_squared_error',
        cv=10).mean()
    cv_errors.append(mse)

# Show Degrees
d_errors = list(zip(degrees, cv_errors))
print(f"cv_errors: {d_errors}")

# Find optimal degree
best_degree = degrees[np.argmin(cv_errors)]
print(f"Best degree by cross-validation: {best_degree}")
```

```
cv_errors: [(1, 1675.6742498921044), (2, 1600.326984342545), (3,
1595.3115754261844), (4, 1593.9356555748932), (5, 1594.4795519650995), (6,
1594.0782857404324), (7, 1592.997234853729), (8, 1592.824347496564), (9,
1595.9062628845313), (10, 1603.793938030312)]
Best degree by cross-validation: 8
```

```
[5]: # plot results
plt.figure(figsize=(8, 4))
plt.plot(degrees, cv_errors, marker='o')
plt.xlabel('Polynomial Degree')
plt.ylabel('10-fold CV Error (MSE)')
plt.title('CV Error for Polynomial Degrees')
plt.grid(True)
plt.show()
```



```
[12]: import statsmodels.api as sm
from ISLP.models import (summarize, poly, ModelSpec as MS)
from statsmodels.stats.anova import anova_lm

# Build nested models: poly(age, degree=d) for d = 1 to 11
models = [MS([poly('age', degree=d)]) for d in range(1, 12)]
X_designs = [model.fit_transform(Wage) for model in models]
fits = [sm.OLS(y, X_).fit() for X_ in X_designs]

# Compare with ANOVA
anova_results = anova_lm(*fits)
print(anova_results)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	2998.0	5.022216e+06	0.0	NaN	NaN	NaN
1	2997.0	4.793430e+06	1.0	228786.010128	143.718829	2.225306e-32
2	2996.0	4.777674e+06	1.0	15755.693664	9.897414	1.671346e-03
3	2995.0	4.771604e+06	1.0	6070.152124	3.813149	5.094478e-02
4	2994.0	4.770322e+06	1.0	1282.563017	0.805681	3.694725e-01
5	2993.0	4.766389e+06	1.0	3932.257665	2.470166	1.161314e-01
6	2992.0	4.763834e+06	1.0	2555.281281	1.605177	2.052697e-01
7	2991.0	4.763707e+06	1.0	126.668985	0.079571	7.778992e-01
8	2990.0	4.756703e+06	1.0	7004.317139	4.399973	3.602325e-02
9	2989.0	4.756701e+06	1.0	2.637537	0.001657	9.675343e-01
10	2988.0	4.756597e+06	1.0	103.201509	0.064829	7.990375e-01

```

[13]: import numpy as np
import matplotlib.pyplot as plt
from ISLP.models import ModelSpec as MS, poly

# Define degrees
degree_cv = 8
degree_anova = 3

# Create ModelSpecs
ms_cv = MS([poly('age', degree=degree_cv)])
ms_anova = MS([poly('age', degree=degree_anova)])

# Fit design matrices
X_cv = ms_cv.fit_transform(Wage)
X_anova = ms_anova.fit_transform(Wage)
y = Wage['wage']


# Fit models
fit_cv = sm.OLS(y, X_cv).fit()
fit_anova = sm.OLS(y, X_anova).fit()

# Generate grid of age values
age_grid = np.linspace(Wage['age'].min(), Wage['age'].max(), 100)
grid_df = pd.DataFrame({'age': age_grid})

# Predict using each model
X_grid_cv = ms_cv.transform(grid_df)
X_grid_anova = ms_anova.transform(grid_df)
pred_cv = fit_cv.predict(X_grid_cv)
pred_anova = fit_anova.predict(X_grid_anova)

# Plot
fig, ax = plt.subplots(figsize=(10, 6))
ax.scatter(Wage['age'], y, color='lightgray', edgecolor='k', alpha=0.6,
           label='Observed Data')
ax.plot(age_grid, pred_anova, label=f'ANOVA Fit (Degree {degree_anova})',
       color='green', linewidth=2)
ax.plot(age_grid, pred_cv, label=f'CV Fit (Degree {degree_cv})', color='blue',
       linewidth=2, linestyle='--')
ax.set_xlabel('Age')
ax.set_ylabel('Wage')
ax.set_title('Polynomial Regression Fits: CV vs ANOVA')
ax.legend()
ax.grid(True)
plt.show()

```



### 1.3 Question 6b.

Fit a step function to predict wage using age, and perform cross-validation to choose the optimal number of cuts. Make a plot of the fit obtained.

```
[15]: import numpy as np
import pandas as pd
import statsmodels.api as sm
from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LinearRegression
from matplotlib.pyplot import subplots
from ISLP import load_data

# Load data again
Wage = load_data('Wage')
X_age = Wage['age']
y = Wage['wage']

# Try step functions with 2 to 10 bins
cv_errors = []
cuts_range = range(2, 11)

for k in cuts_range:
    # Bin ages into k intervals
```

```

Wage['age_bin'] = pd.cut(X_age, bins=k)

# One-hot encode the age bins (drop first to avoid collinearity)
X_dummies = pd.get_dummies(Wage['age_bin'], drop_first=True)

# Cross-validation using sklearn
model = LinearRegression()
scores = cross_val_score(model, X_dummies, y, cv=10,
↪scoring='neg_mean_squared_error')
cv_errors.append(-scores.mean())

# Find best number of bins
best_k = cuts_range[np.argmin(cv_errors)]
print(f"cv_errors:{cv_errors}")
print(f"Best number of age bins (cuts): {best_k}")

```


cv\_errors: [1733.6247654372023, 1682.924818990572, 1635.2904177897497,  
1631.8552686011149, 1623.656547415575, 1611.4169308778028, 1600.0110941892897,  
1611.1198831331985, 1605.6087735908577]  
Best number of age bins (cuts): 8

```

[16]: import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))
plt.plot(cuts_range, cv_errors, marker='o')
plt.xlabel('Number of Age Cuts (Bins)')
plt.ylabel('10-fold CV Error (MSE)')
plt.title('Step Function CV Error vs. Number of Cuts')
plt.grid(True)
plt.show()

```



```
[17]: # Recreate best step function
Wage['age_bin'] = pd.cut(X_age, bins=best_k)
X_best = pd.get_dummies(Wage['age_bin'], drop_first=True)
model = LinearRegression().fit(X_best, y)

# Predict over age grid using bin mapping
age_grid = np.linspace(X_age.min(), X_age.max(), 1000)
age_df = pd.DataFrame({'age': age_grid})
age_df['age_bin'] = pd.cut(age_df['age'], bins=best_k)
X_grid = pd.get_dummies(age_df['age_bin'], drop_first=True)

# Align dummy columns in case of mismatch
X_grid = X_grid.reindex(columns=X_best.columns, fill_value=0)


# Predict wages
wage_pred = model.predict(X_grid)

# Plot
fig, ax = subplots(figsize=(10, 6))
ax.scatter(X_age, y, color='lightgray', edgecolor='k', alpha=0.5,
           label='Observed Data')
ax.plot(age_grid, wage_pred, color='red', linewidth=2, label=f'Step Function
           (k={best_k})')
```

```

ax.set_xlabel('Age')
ax.set_ylabel('Wage')
ax.set_title('Step Function Fit for Wage vs. Age')
ax.legend()
ax.grid(True)
plt.show()

```



#### 1.4 Question 10.

This question relates to the College data set.

```

[24]: from ISLP import load_data

College = load_data('College')
College.head()

```

```

[24]: Private  Apps  Accept  Enroll  Top10perc  Top25perc  F.Undergrad  \
0    Yes  1660   1232    721         23         52        2885
1    Yes  2186   1924    512         16         29        2683
2    Yes  1428   1097    336         22         50        1036
3    Yes   417    349    137         60         89         510
4    Yes   193   146     55         16         44         249

P.Undergrad  Outstate  Room.Board  Books  Personal  PhD  Terminal  \

```



0	537	7440	3300	450	2200	70	78
1	1227	12280	6450	750	1500	29	30
2	99	11250	3750	400	1165	53	66
3	63	12960	5450	450	875	92	97
4	869	7560	4120	800	1500	76	72

	S.F.Ratio	perc.alumni	Expend	Grad.Rate
0	18.1	12	7041	60
1	12.2	16	10527	56
2	12.9	30	8735	54
3	7.7	37	19016	59
4	11.9	2	10922	15

### 1.5 Question 10a.

Split the data into a training set and a test set. Using out-of-state tuition as the response and the other variables as the predictors perform forward stepwise selection on the training set in order to identify a satisfactory model that uses just a subset of the predictors.

```
[ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
from ISLP import load_data

# 1) Load the College data
College = load_data('College')

# 2) Separate response and predictors
y = College['Outstate'].astype(float)          # ensure numeric
X = College.drop(columns='Outstate')

# 3) One-hot encode any categoricals, then cast ALL columns to float
X = pd.get_dummies(X, drop_first=True).astype(float)

# 4) Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=0
)

# 5) Forward stepwise selection (AIC-based)
def forward_selection(X, y):
    remaining = set(X.columns)
    selected = []
    current_aic = np.inf

    while remaining:
```

```

        best_candidate = None
        best_aic = current_aic
        for cand in remaining:
            Xc = sm.add_constant(X[selected + [cand]], has_constant='add')
            model = sm.OLS(y, Xc).fit()
            if model.aic < best_aic:
                best_aic = model.aic
                best_candidate = cand

        # If we found an improvement, keep it
        if best_candidate is not None:
            remaining.remove(best_candidate)
            selected.append(best_candidate)
            current_aic = best_aic
        else:
            break

    return selected, current_aic

# 6) Run forward selection on the training set
selected_vars, final_aic = forward_selection(X_train, y_train)

# 7) Report the result
print("Forward stepwise selected predictors:")
for v in selected_vars:
    print("  -", v)
print(f"Final training-set AIC: {final_aic:.2f}")

### I VIBE CODED THIS... THE BOOKS EXAMPLE SUCKS! HOURS OF MY LIFE WASTED!!!

```

Forward stepwise selected predictors:

- Expend
- Private\_Yes
- Room.Board
- perc.alumni
- Terminal
- Grad.Rate
- Top10perc
- Personal
- Accept
- F.Undergrad
- Enroll
- Apps
- Top25perc

Final training-set AIC: 9765.02

## 1.6 Question 10b.

Fit a GAM on the training data, using out-of-state tuition as the response and the features selected in the previous step as the predictors. Plot the results, and explain your findings.

```
[45]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from pygam import LinearGAM, s as s_gam
from matplotlib.pyplot import subplots

# 1) Load & preprocess (as in 10a):
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)

# 2) The forward-stepwise selected predictors:
selected_vars = [
    'Expend', 'Private_Yes', 'Room.Board', 'perc.alumni', 'Terminal',
    'Grad.Rate', 'Top10perc', 'Personal', 'Accept',
    'F.Undergrad', 'Enroll', 'Apps', 'Top25perc'
]
Xg_train = X_train[selected_vars].values

# 3) Build the GAM term list correctly:
terms = s_gam(0)
for i in range(1, Xg_train.shape[1]):
    terms = terms + s_gam(i)














# 4) Fit the GAM:
gam = LinearGAM(terms)
gam.gridsearch(Xg_train, y_train.values)

# 5) Plot each smooth effect (no CIs)
fig, axes = subplots(len(selected_vars), 1,
                      figsize=(8, 4 * len(selected_vars)))
for i, var in enumerate(selected_vars):
    # grid for partial dependence
    XX = gam.generate_X_grid(term=i)
    pdep = gam.partial_dependence(term=i, X=XX)

    ax = axes[i]
    ax.plot(XX[:, i], pdep, linewidth=2)
    ax.scatter(Xg_train[:, i], y_train,
               facecolor='lightgray', alpha=0.3, s=20)
    ax.set_xlabel(var)
    ax.set_ylabel(f'f({var})')
    ax.set_title(f'Smooth effect of {var}')
    ax.grid(True)
```

```
fig.tight_layout()
```

```
0% (0 of 11) | Elapsed Time: 0:00:00 ETA:  --:--:--
9% (1 of 11) |## Elapsed Time: 0:00:00 ETA:  0:00:02
18% (2 of 11) |#### Elapsed Time: 0:00:00 ETA:  0:00:02
27% (3 of 11) |##### Elapsed Time: 0:00:00 ETA:  0:00:01
36% (4 of 11) |##### Elapsed Time: 0:00:01 ETA:  0:00:01
45% (5 of 11) |##### Elapsed Time: 0:00:01 ETA:  0:00:01
54% (6 of 11) |##### Elapsed Time: 0:00:01 ETA:  0:00:01
63% (7 of 11) |##### Elapsed Time: 0:00:02 ETA:  0:00:01
72% (8 of 11) |##### Elapsed Time: 0:00:02 ETA:  0:00:01
81% (9 of 11) |##### Elapsed Time: 0:00:03 ETA:  0:00:00
90% (10 of 11) |##### Elapsed Time: 0:00:03 ETA:  0:00:00
100% (11 of 11) |##### Elapsed Time: 0:00:04 Time:  0:00:04
```



### 1.7 Question 10c.

Evaluate the model obtained on the test set, and explain the results obtained.

```
[46]: from sklearn.metrics import mean_squared_error, r2_score

# Prepare the test-set matrix
Xg_test = X_test[selected_vars].values

# Predict out-of-state tuition on the test set
y_pred = gam.predict(Xg_test)

# Compute test-set MSE and  $R^2$ 
mse_test = mean_squared_error(y_test, y_pred)
r2_test = r2_score(y_test, y_pred)

print(f"Test   MSE: {mse_test:.2f}")
print(f"Test    $R^2$  : {r2_test:.3f}")
```

```
Test   MSE: 3378526.25
```

```
Test    $R^2$  : 0.800
```

Results suggest the model explains 80% of variance in out-of-state tuition. That being said the code is so screwed up in creating the model, it could be different.

### 1.8 Question 10d.

For which variables, if any, is there evidence of a non-linear relationship with the response?

Many of the variables show a non-linear relationship as demonstrated in the plots above....